

```

# Import required libraries
import pandas as pd
import numpy as np
import time
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris, fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, mean_squared_error, confusion_matrix
from sklearn.linear_model import LinearRegression
from sklearn.cluster import KMeans
from mlxtend.frequent_patterns import apriori, association_rules
from mlxtend.preprocessing import TransactionEncoder
print(" ♦ Task 1: Comparative Study of Data Mining Techniques\n")
# -----
# 1. Classification - Decision Tree
# -----
iris = load_iris()
X = iris.data
y = iris.target
X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(X, y, test_size=0.3, random_state=42)
start_time = time.time()
clf = DecisionTreeClassifier()
clf.fit(X_train_c, y_train_c)
y_pred_c = clf.predict(X_test_c)
end_time = time.time()
print(" ♦ Classification (Decision Tree):")
print("Accuracy:", accuracy_score(y_test_c, y_pred_c))
print("Precision:", precision_score(y_test_c, y_pred_c, average='macro'))
print("Recall:", recall_score(y_test_c, y_pred_c, average='macro'))
print("Execution Time:", round(end_time - start_time, 5), "seconds\n")
# ♦ Visualization: Confusion Matrix
plt.figure(figsize=(5, 4))
sns.heatmap(confusion_matrix(y_test_c, y_pred_c), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Decision Tree Classifier")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
# -----
# 2. Regression - California Housing
# -----
california = fetch_california_housing()
X_cal = pd.DataFrame(california.data, columns=california.feature_names)
y_cal = california.target

X_train_r, X_test_r, y_train_r, y_test_r = train_test_split(X_cal, y_cal, test_size=0.3, random_state=42)
start_time = time.time()
reg = LinearRegression()
reg.fit(X_train_r, y_train_r)
y_pred_r = reg.predict(X_test_r)
end_time = time.time()
print(" ♦ Regression (Linear Regression on California Housing):")
print("Mean Squared Error:", mean_squared_error(y_test_r, y_pred_r))
print("R² Score:", reg.score(X_test_r, y_test_r))
print("Execution Time:", round(end_time - start_time, 5), "seconds\n")
# ♦ Visualization: Actual vs Predicted
plt.figure(figsize=(6, 5))
sns.scatterplot(x=y_test_r, y=y_pred_r, alpha=0.5)
plt.plot([min(y_test_r), max(y_test_r)], [min(y_test_r), max(y_test_r)], 'r--')
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted Prices (Linear Regression)")
plt.show()
# -----

```

```

# 3. Clustering - KMeans
# -----
start_time = time.time()
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
kmeans.fit(X)
labels = kmeans.labels_
end_time = time.time()
print(" ♦ Clustering (KMeans on Iris Data):")
print("Cluster Centers:\n", kmeans.cluster_centers_)
print("Execution Time:", round(end_time - start_time, 5), "seconds\n")
# ♦ Visualization: Clusters (first two features)
plt.figure(figsize=(6, 5))
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', edgecolor='k', s=50)
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], s=200, c='red', marker='X')
plt.xlabel(iris.feature_names[0])
plt.ylabel(iris.feature_names[1])
plt.title("KMeans Clustering on Iris Dataset")
plt.show()
# -----
# 4. Association Rule Mining - Apriori
# -----
dataset = [
    ['Milk', 'Bread', 'Butter'],
    ['Bread', 'Eggs'],
    ['Milk', 'Bread', 'Butter', 'Eggs'],
    ['Bread', 'Butter'],
    ['Milk', 'Butter'],
    ['Milk', 'Bread'],
    ['Butter', 'Eggs']
]
te = TransactionEncoder()
te_data = te.fit(dataset).transform(dataset)
df = pd.DataFrame(te_data, columns=te.columns_)
start_time = time.time()
frequent_itemsets = apriori(df, min_support=0.3, use_colnames=True)
rules = association_rules(frequent_itemsets, metric='confidence', min_threshold=0.6)
end_time = time.time()
print(" ♦ Association Rule Mining (Apriori):")
print("Frequent Itemsets:\n", frequent_itemsets)
print("\nGenerated Rules:\n", rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
print("Execution Time:", round(end_time - start_time, 5), "seconds")
# ♦ Visualization: Support vs Confidence
plt.figure(figsize=(6, 5))
sns.scatterplot(x=rules['support'], y=rules['confidence'], size=rules['lift'], hue=rules['lift'], palette='co
plt.title("Support vs Confidence (Association Rules)")
plt.xlabel("Support")
plt.ylabel("Confidence")
plt.grid(True)
plt.show()

```



## ♦ Task 1: Comparative Study of Data Mining Techniques


## ♦ Classification (Decision Tree):

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

Execution Time: 0.01004 seconds




## ♦ Regression (Linear Regression on California Housing):

Mean Squared Error: 0.5305677824766758

 $R^2$  Score: 0.595770232606166

Execution Time: 0.07853 seconds



## ♦ Clustering (KMeans on Iris Data):

Cluster Centers:

[[5.9016129 2.7483871 4.39354839 1.43387097]


[5.006 3.428 1.462 0.246 ]

[6.85 3.07368421 5.74210526 2.07105263]]

Execution Time: 0.41004 seconds

## KMeans Clustering on Iris Dataset

4.5



#### ♦ Association Rule Mining (Apriori):

Frequent Itemsets:

	support	itemsets
0	0.714286	(Bread)
1	0.714286	(Butter)
2	0.428571	(Eggs)
3	0.571429	(Milk)
4	0.428571	(Butter, Bread)
5	0.428571	(Milk, Bread)
6	0.428571	(Milk, Butter)

Generated Rules:

	antecedents	consequents	support	confidence	lift
0	(Butter)	(Bread)	0.428571	0.60	0.84
1	(Bread)	(Butter)	0.428571	0.60	0.84
2	(Milk)	(Bread)	0.428571	0.75	1.05
3	(Bread)	(Milk)	0.428571	0.60	1.05
4	(Milk)	(Butter)	0.428571	0.75	1.05
5	(Butter)	(Milk)	0.428571	0.60	1.05

Execution Time: 0.02534 seconds

Support vs Confidence (Association Rules)

